School of Computer Science and Engineering
Faculty of Engineering
UNSW Sydney

UltraCast

BY

Daniel Latimer   **Connor O'Shea**   **Kevin Chan**   **Oliver Richards**   **Peter Kerr**
   z5115175   **z5115177**   **z5113136**   **z5157383**   **z5115807**

Submitted: November 14, 2020
Tutor: Tatjana Zrimec

Contents

1 Overview

1.1 Introduction

With over 100 million monthly listeners[4] and a steadily increasing user base, there is no doubt that podcasts are a greatly enriching source of information and entertainment for a large variety of individuals.

Although they are highly valuable, it can be difficult to find podcasts that are of interest to a particular user amidst the 1 million[4] that are already available. Thus, podcast streaming services (such as UltraCast) have been created, to provide a centralised place for exploring and discovering new podcasts that are valuable to the listener.

However, all of the web based podcast streaming services available lack many important features, and their interfaces leave much to be desired. For example, there is no streaming service that allows the user to bookmark certain parts of a podcast, nor take notes at certain timestamps. It is even difficult to find a service that allows the listener to change the playback speed of the podcast.

UltraCast combines all of the most important features together into a single package with a web-based podcast streaming service.

UltraCast differentiates itself from competitors by allowing users to:

- Follow friends to see what they have been listening to

- Create *Streams* of podcasts to find interesting podcasts

- Create bookmarks inside podcast episodes

- Monitor episode and podcast play metrics

1.2 Project Requirements

The minimun project requirements from the specifications are:

- Listeners must be able to search for podcasts that interest them by keywords, resulting in a list of matching podcast titles, where the total number of subscriptions on the UltraCast platform (function described later) for each podcast is shown next to the title

- Listeners must be able to select a podcast show from returned search results to view its full details, including its title, description, any author details that exist, as well as a list of episodes for the show

- Listeners must be able to play a selected episode within a podcast show, and once that episode starts being played, the listener must be able to also clearly see this episode marked as "Played"

- Listeners must be able to subscribe or unsubscribe from a podcast show Listeners must be able to see the latest episode available for each show that they subscribed to in a "Podcast Subscription Preview" panel

- Listeners must be notified by the platform when a new episode for a show they are subscribed appears

- Listeners must be able to see a history of the podcast episodes that they have played, sorted in order from most recently played to least recently played

- UltraCast must be able to recommend new podcast shows to a listener based on at least information about the podcast shows they are subscribed to, podcast episodes they have recently played, and their past podcast searches

The following additional requirements have also been implemented:

- Listeners should be able to follow their *friends* and view the podcasts that their *friends* have recently listened to

- Listeners should be able to create *Streams* based off search queries that they can use to find interesting podcasts

- Listeners should be able to add bookmarks with a name and description to podcast episodes as they listen to them

- Content creators should be able to create and upload podcasts and podcast episodes

- Content creators should be able to monitor analytics of their uploaded podcasts related to their listeners

1.3 System Architecture

The highlevel system architecture can be seen in Figure 1. It can be seen that our end users will be podcast listeners and content creators who connect to the presentation layer which is powered by a ReactJS application. ReactJS was selected as the framework for the frontend application due to .... TODO justification of React

Flask, a micro-framework, [1] was used for the web-server, due ease of use which allowed for rapid development. The React application communicates with Flask web-server through a GraphQL API: a scalable alternative to the popular REST API[2]. The Flask web-server also contains a recommendation service which drives recommendation functionality described in TODO. The Flask web-server uploads static files (episode audio and podcast covers) to a static file storage server while the urls for these static files along with all other data is stored in MongoDB. MongoDB, a NoSQL database, was selected to store metadata due to its scalability [3].
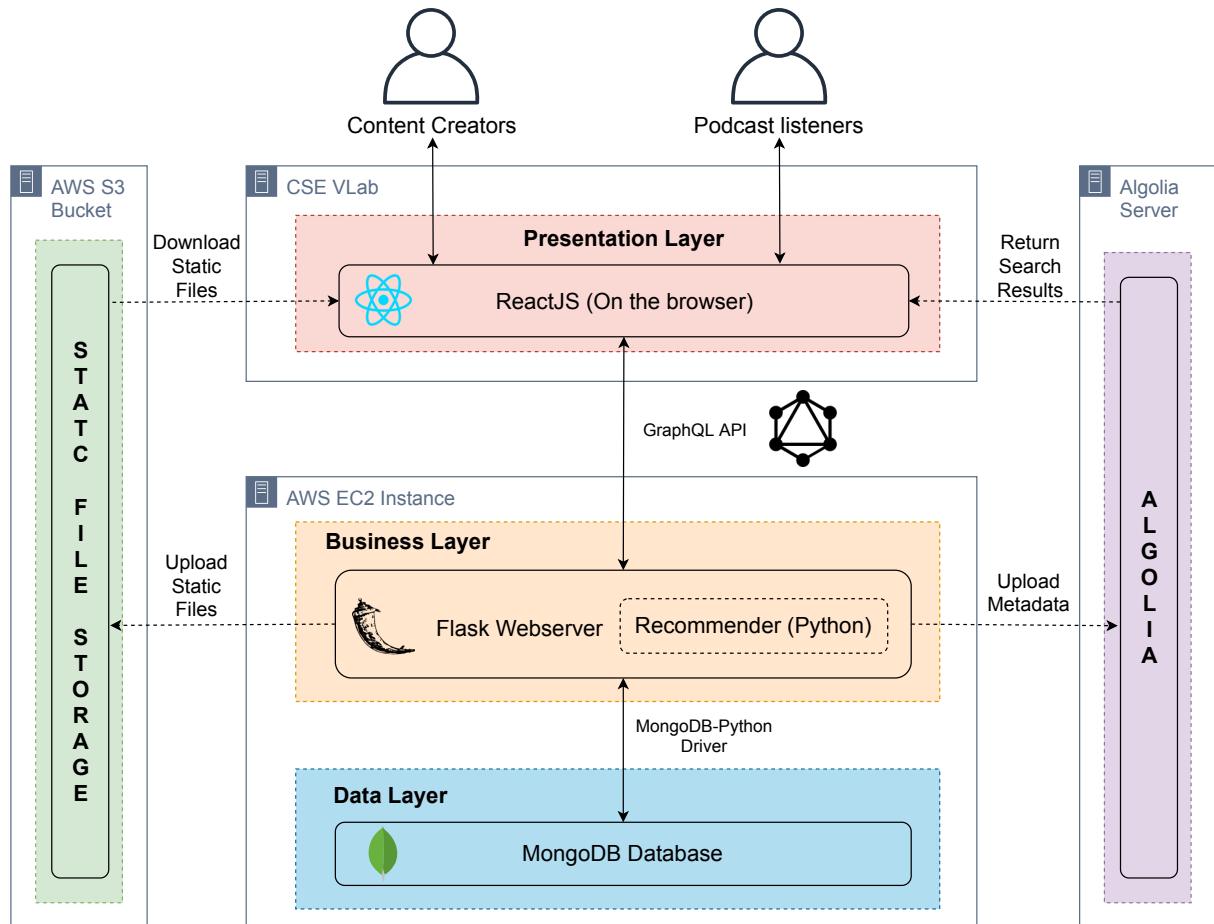

Figure 1: UltraCast System Architecture

A static file storage server (AWS S3 Bucket) was used instead of MongoDB due to the lower costs and increased performance associated with storing larger files, particularly audio files TODO-cite. MongoDB was not supported by Debian 6 (the Linux environment on the VLab machine), so a remote server (AWS EC2 instance) was used to host MongoDB. To improve query performance, the entire business layer was migrated to the remote server. (TODO - should probs cite why query performance improved)

2 Functionalities and Implementation Challenges

2.1 Functionalities

2.1.1 Viewing and Searching

TODO reference and write description

Story ID	Functionality	Project Objective
UL-2 UL-3	Use keywords to search for podcasts, return list of podcasts (See UL-4 for format) View the total number of subscribers for each podcast returned from a search	Listeners must be able to search for podcasts that interest them by keywords, resulting in a list of matching podcast titles, where the total number of subscriptions on the UltraCast platform (function described later) for each podcast is shown next to the title
UL-4	View the title, description, author details and list of episodes for a podcast	Listeners must be able to select a podcast show from returned search results to view its full details, including its title, description, any author details that exist, as well as a list of episodes for the show
UL-14	Login as specific user	-
UL-24	View a title, length, upload date for episodes	-
UL-29	Save search as a "Stream"	-
UL-41	Signup as user	-

2.1.2 Playing Podcast Episodes

TODO reference and write description

Story ID	Functionality	Project Objective
UL-5 UL-6	Play episodes Once episode starts being played it is marked as played	Listeners must be able to play a selected episode within a podcast show, and once that episode starts being played, the listener must be able to also clearly see this episode marked as "Played"
UL-18	Pause episode that is playing	-
UL-19	Adjust playback volume	-
UL-20	Skip to next episode, previous episode and start of current episode	-
UL-21	Jump to a point in an episode	-
UL-22	Adjust playback speed	-

Story ID	Functionality	Project Objective
UL-23	Auto-play episodes in a podcast (after added to playlist)	-
UL-26	"Bookmark" a point in an episode with a title and description	-

2.1.3   Recommendation and Following

TODO reference and write description

Story ID	Functionality	Project Objective
UL-10	View episode history	Listeners must be able to see a history of the podcast episodes that they have played, sorted in order from most recently played to least recently played
UL-11	Episode history is sorted by most recent to least recent	
UL-12	Podcast recommendations are based on: Existing subscriptions recently played episodes and past searches	UltraCast must be able to recommend new podcast shows to a listener based on at least information about the podcast shows they are subscribed to, podcast episodes they have recently played, and their past podcast searches
UL-13	A "recommended" panel shows recommended podcasts	
UL-18	Follow users, view their listen history	-

2.1.4   Creator Mode

TODO reference and write description

Story ID	Functionality	Project Objective
UL-15	Create podcasts and episodes	-
UL-16	Delete podcasts and episodes	
UL-17	Update podcasts and episodes	-
UL-27	Access to podcast and episode viewer metrics	-

2.1.5   Subscribing

TODO reference and write description

Story ID	Functionality	Project Objective
UL-7	Subscribe to podcasts	Listeners must be able to see a history of the podcast episodes that they have played, sorted in order from most
UL-8	Unsubscribe to podcasts	recently played to least recently played

UL-9	User receives notification for each new episode in a podcast they are subscribed to	Listeners must be notified by the platform when a new episode for a show they are subscribed appear
UL-30	The latest episode for each subscribed podcast is linked in the "Subscriptions" page	Listeners must be able to see the latest episode available for each show that they subscribed to in a "Podcast Subscription Preview" panel

2.2 System

2.3 Implementation Challenges

2.3.1 Backend Stack

The backend of UltraCast employs an unusual technology stack, with MongoDB as a persistence layer, flask as a webserver framework and graphql (via graphene and graphene-mongo libraries) as an API layer. This created difficulties in implementing common web-app functionalities due to (1) a lack of documentation on the libraries being used and (2) no online examples implementing these functionalities with this stack.

User Authentication Implementing user authentication for the backend was a non-trivial task because the Graphene and Graphene-Mongo libraries which are used for the API layer do not natively support this functionality. A major challenge in applying general purpose authentication libraries, for example flask-jwt[1], is that only one route is used for all API calls. Some of these API calls need to be authenticated e.g. deleting a podcast where others should not e.g. signing up to the site. The Flask-GraphQL-Auth library[2] provides the required authentication methods, however, it is not actively maintained. After much research, user authentication was implemented using the flask-jwt-extended library[3]. This library allows authentication to be required on a per-function level, rather than for an entire route. Hence, certain mutations and queries can be protected with user authentication where required. The frontend calls a signin mutation which returns a Json Web Token (JWT). This mutation does not require authentication. The frontend then stores this JWT as a cookie and sends it in the header of any future GraphQL API requests.

Resolving Nested Queries While testing the frontend, it was discovered that some backend GraphQL queries were taking upwards of one minute to return. The site was still responsive, however it took a long time for recommended podcasts to be displayed. Further investigation revealed that where nested references were used in the database models, and the GraphQL query involved dereferencing these references, the Graphene-Mongo library would perform one database operation per parent node. These database operations are performed sequentially. Since the MongoDB instance is hosted in the cloud, each database operation takes some number of milliseconds due to network latency. When a large number of parent nodes were fetched, this resulted in very slow queries. It was not

---

[1]Available at https://github.com/mattupstate/flask-jwt
[2]Available at https://github.com/NovemberOscar/Flask-GraphQL-Auth
[3]Available at https://github.com/vimalloc/flask-jwt-extended

feasible to modify the Graphene-Mongo libary to issue less database operations. Hence, the decision was made to move the GraphQL API webserver to the same cloud container as the MongoDB instance. This improved the time for some queries from over fourty seconds to less than a second.

Database Integrity

Populating the Site   To build a meaningful recommendation system, the website must have a reasonable amount of podcasts already uploaded to it. Since UltraCast has not been released, there are no users to generate this data. To allow for experimentation with different approaches to recommending podcasts to users, a podcast dataset was scraped. It was difficult to find a suitable dataset that contained the required category, sub-category and keyword tags for podcasts that did not impose commercial obligations on UltraCast (due to terms of use of the dataset). A dataset which is an aggregation of public domain podcasts was found and scraped, providing over 200 podcasts and 2000 podcast episodes for the site.

3   User Manual

3.1   Software Setup Instructions

3.2   Site Usage and Functionality Guide

References

[1] Flask: web development, one drop at a time. `https://flask.palletsprojects.com/en/1.1.x/`. Accessed: 2020-11-09.

[2] Graphql. `https://graphql.org/`. Accessed: 2020-11-09.

[3] MongoDB. The most popular database for modern apps — MongoDB, 2020. `http://mongodb.com`, Last accessed on 2020-11-08.

[4] Gavin Whitner. Podcast statistics (2020). `https://musicoomph.com/podcast-statistics/`, Sep 2020. Accessed: 2020-11-08.